

# Introduction to Delta Lake

By Gerhard Brueckl



SSAS  
MAESTRO  
by Microsoft



@gbrueckl



blog.gbrueckl.at



[gerhard@gbrueckl.at](mailto:gerhard@gbrueckl.at)



<https://github.com/gbrueckl>



[DatabricksPS](#)



[Databricks VSCode](#)



[PowerBI Connector](#)



[www.paiqo.com](http://www.paiqo.com)



# Agenda

- What is Delta Lake?
- How does it work?
- What do I need to watch out for?

## Delta Lake – [delta.io](https://delta.io)

Delta Lake is an open-source storage layer that brings ACID transactions to Apache Spark™ and big data workloads.

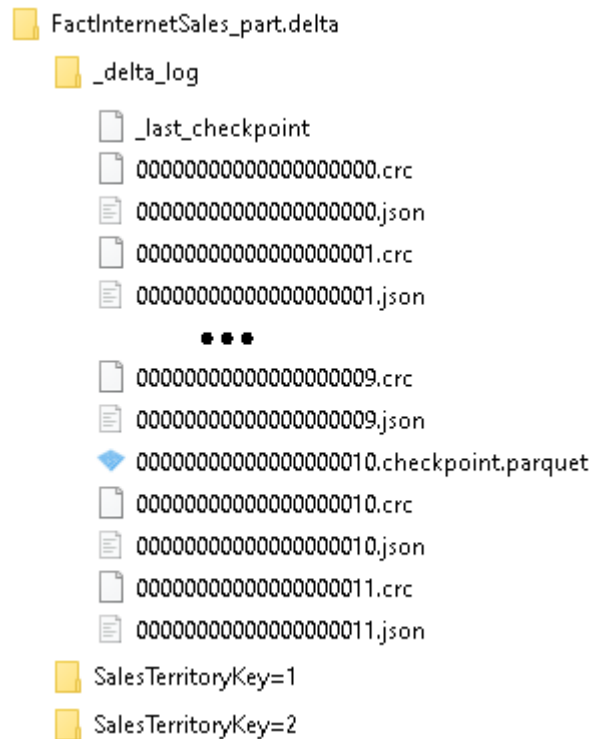
- ACID compliant transactions
  - Optimistic Concurrency Control
- Support for UPDATE / MERGE
- Time-Travel
- Schema enforcement and evolution
  - Across multiple files/folders
- Batch & Streaming
- 100% compatible with Apache Spark

# Delta Lake – [delta.io](https://delta.io)

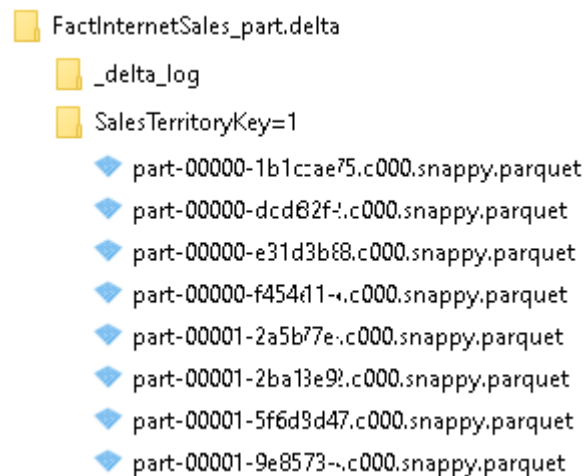
- Self-contained table
  - Meta-data / schema
  - Partitions
  - Data
- Can reside on any storage
- Based on Apache Parquet and JSON
- Transaction isolation on partition level
- Native support in Databricks / Spark (Azure Synapse, ...)
- Restore of previous version, cloning of tables
- APIs for SQL, Python, Scala, ...

# Delta Lake – [delta.io](https://delta.io)

## \_delta\_log



## data



# DDL Operations – CREATE TABLE AS SELECT

```
CREATE TABLE IF NOT EXISTS DimProduct
USING DELTA
PARTITIONED BY (ProductSubcategoryKey)
LOCATION '/mnt/adls/tables/DimProduct'
TBLPROPERTIES ('myKey' = 'myValue')
AS SELECT *
FROM stage.Products
```

- Column definitions are taken from **SELECT**

# DDL Operations – CREATE TABLE (existing)

```
CREATE TABLE IF NOT EXISTS DimProduct  
USING DELTA  
LOCATION '/mnt/adls/tables/DimProduct'
```

- Avoid defining any additional information – this is all handled by the delta log!
  - Everything within the delta-folder is self-contained!
  - If you still define it, it has to match the existing information in the delta log



# DML Operations – Delta Lake - MERGE

```
MERGE INTO targetTable
USING updatesTable
    ON targetTable.year = 2021
    AND targetTable.id = updatesTable.id
WHEN MATCHED THEN
    UPDATE SET targetTable.value = updatesTable.value
WHEN NOT MATCHED THEN
    INSERT (id, year, value)
    VALUES (updatesTable.id, updatesTable.year, updatesTable.value)
```

- UPDATE, INSERT and DELETE
  - in one statement
  - In one transaction

# DML Operations – Delta Lake - UPDATE

User

Product	Price
Notebook	900 €
PC	1,500 €
Tablet	500 €

```
UPDATE TABLE DimProduct
SET Price = 1300
WHERE Product = 'PC'
```

Product	Price
Notebook	900 €
PC	1,300 €
Tablet	500 €

\_delta\_log

```
000000000.json
"add": {
  "path": "part-01.parquet",
  ...
}
```

```
000000001.json
"remove": { "path": "part-01.parquet", ... },
"add": { "path": "part-02.parquet", ... }
```

Storage



Parquet  
part-01  
(3 rows)



Parquet  
part-01  
(3 rows)



Parquet  
part-02  
(3 rows)

# DML Operations - Delta Lake – DELETE

User

Product	Price
Notebook	900 €
PC	1,300 €
Tablet	500 €

```
DELETE FROM DimProduct
WHERE Product = 'PC'
```

Product	Price
Notebook	900 €
Tablet	500 €

\_delta\_log

000000000.json

```
"add": {
  "path": "part-01.parquet",
  ... }
```

000000001.json

```
"remove": {
  "path": "part-01.parquet", ... },
"add": {
  "path": "part-02.parquet", ... }
```

000000002.json

```
"remove": { "path": "part-02.parquet", ... },
"add": { "path": "part-03.parquet", ... }
```

Storage



Parquet  
part-01  
(3 rows)



Parquet  
part-02  
(3 rows)



Parquet  
part-01  
(3 rows)



Parquet  
part-02  
(3 rows)



Parquet  
part-03  
(2 rows)

# DML Operations - Delta Lake – INSERT

User

Product	Price
Notebook	900 €
PC	1,300 €
Tablet	500 €

```
INSERT INTO DimProduct
VALUES ('Monitor', 200)
```

Product	Price
Notebook	900 €
Tablet	500 €
Monitor	200 €

\_delta\_log

000000000.json

000000001.json

```
"remove": {
  "path": "part-01.parquet", ... },
"add": {
  "path": "part-02.parquet", ... }
```

000000002.json

```
"remove": { "path": "part-02.parquet", ... },
"add": { "path": "part-03.parquet", ... }
```

000000003.json

```
"add": { "path": "part-04.parquet", ... }
```

Storage



Parquet  
part-01  
(3 rows)



Parquet  
part-02  
(3 rows)



Parquet  
part-03  
(2 rows)



Parquet  
part-01  
(3 rows)



Parquet  
part-02  
(3 rows)



Parquet  
part-03  
(2 rows)



Parquet  
part-04  
(1 row)

# DML Operations - Delta Lake – VACUUM

User

Product	Price
Notebook	900 €
PC	1,300 €
Tablet	500 €

VACUUM DimProduct

Product	Price
Notebook	900 €
Tablet	500 €
Monitor	200 €

\_delta\_log

000000000.json  
000000001.json

000000002.json  
"remove": {  
  "path": "part-02.parquet", ... },  
"add": {  
  "path": "part-03.parquet", ... }  
000000003.json  
"add": {  
  "path": "part-04.parquet",  
  ... }  
}

Storage




Parquet  
part-01  
(3 rows)




Parquet  
part-02  
(3 rows)



Parquet  
part-03  
(2 rows)



Parquet  
part-03  
(2 rows)



Parquet  
part-04  
(1 row)

# Delta Lake – UPDATE/DELETE/MERGE

- Operations are logged in `_delta_log`
  - Old files are removed/invalidated (logically)
  - New files are added/referenced
- Conflicts have to be handled by the User!
- Always results in new, duplicated files! Even a **DELETE** can!

Can create A LOT of files!

# Delta Lake – Storage Consumption

Each DML Operation can potentially double the storage consumption (if all files are touched)

- UPDATE/DELETE/MERGE
- OPTIMIZE (!)

VACUUM is mandatory !

It makes sense to monitor the table size and storage consumption using a [script](#) and storing the results in a table